

Operating Systems - Processes

Mridul Sankar Barik

Jadavpur University


2024

Slide Credits

- Most of the slides are adapted from the companion lecture slides for the text book by Avi Silberschatz, Peter Baer Galvin, Greg Gagne


Process Concept

- Process is a program in execution
- A process includes
 - **Program code** (text section)
 - **Program counter and other processor registers**
 - Defines the current state of the process
 - **Stack**: contains temporary data
 - Function parameters, return addresses, local variables
 - **Data section**: contains global data
 - **Heap**: contains memory dynamically allocated during run time
- A Program is a **passive** entity stored on disk (executable file), whereas process is an **active** entity




Process States

- **New:** The process is being created
- **Running:** Instructions are being executed
- **Waiting:** The process is waiting for some event to occur
- **Ready:** The process is waiting to be assigned to a processor
- **Terminated:** The process has finished execution

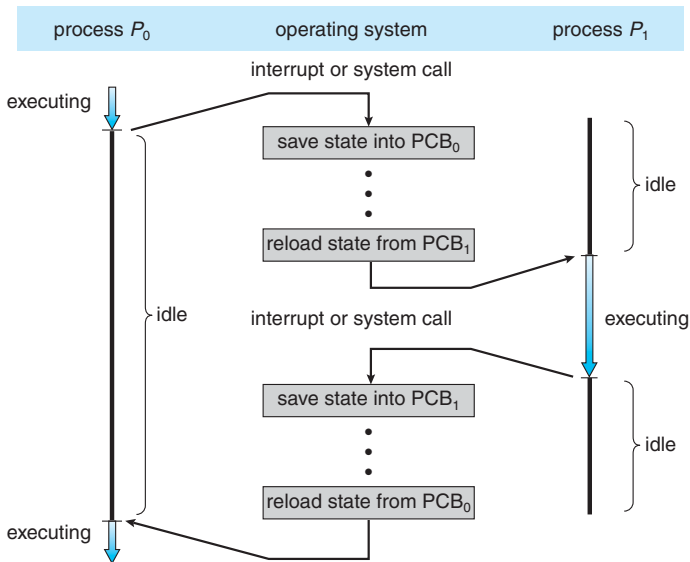


Process Control Block (PCB)

- Each process is represented in the OS by a PCB
 - Process State
 - Program Counter
 - CPU Registers
 - CPU Scheduling Information
 - process priority
 - pointer to scheduling queues etc.
 - Memory Management Information
 - values of base, limit registers
 - page table or segment table etc.
 - Accounting Information
 - amount of CPU and real time used
 - time limits
 - job or process numbers etc.
 - I/O Status Information
 - list of I/O devices allocated to this process
 - list of open files etc.



CPU Switch From Process to Process




Threads

- So far, process has a single thread of execution
- Consider having multiple program counters per process
 - Multiple locations can execute at once
 - Multiple threads of control \Rightarrow threads
- Must then have storage for thread details, multiple program counters in PCB

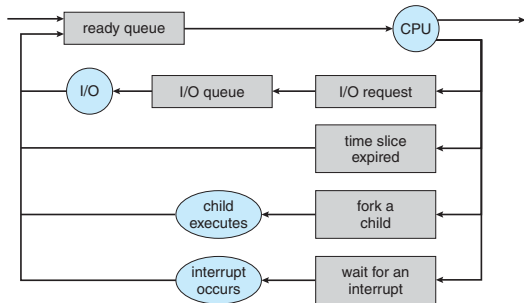
Process Scheduling I

- Maximize CPU use, quickly switch processes onto CPU for time sharing
- **Processes scheduler** selects among available processes for next execution on CPU
- Maintains scheduling queues of processes
 - **Job queue** – set of all processes in the system
 - **Ready queue** – set of all processes residing in main memory, ready and waiting to execute
 - **Device queues** – set of processes waiting for an I/O device
- Processes migrate among the various queues



Process Scheduling II

- A new process is initially put in the ready queue, waits there until it is selected for execution
- Once a process executes it may
 - Issue an I/O request, and be placed in an I/O queue
 - Create a new process and wait for its termination
 - Be removed forcibly from the CPU as a result of an interrupt




Schedulers I

- **Short-term scheduler** (or **CPU scheduler**) – selects which process should be executed next and allocates CPU
 - Sometimes the only scheduler in a system
 - Short-term scheduler is invoked frequently (milliseconds) (must be fast)
- **Long-term scheduler** (or **job scheduler**) – selects which processes should be brought into the ready queue
 - Long-term scheduler is invoked infrequently (seconds, minutes) (may be slow)
 - The long-term scheduler controls the **degree of multiprogramming**
- Processes can be described as either:
 - **I/O-bound process** – spends more time doing I/O than computations, many short CPU bursts
 - **CPU-bound process** – spends more time doing computations; few very long CPU bursts
- Long-term scheduler strives for good process mix

Schedulers II

- **Medium-term scheduler** can be added if degree of multiple programming needs to decrease
 - Remove process from memory, store on disk, bring back in from disk to continue execution: **swapping**



Context Switch

- When CPU switches to another process, the system must **save the state** of the old process and load the **saved state** for the new process via a **context switch**
- **Context** of a process is represented in the PCB of a process
- **Context-switch time** is pure overhead
 - The system does no useful work while switching
 - The more complex the OS and the PCB \Rightarrow the longer the context switch
- Context switch time is dependent on hardware support
 - Sun UltraSPARC provides multiple sets of registers
 - Context switch simply requires changing the pointer to the current register set