

Advanced Operating Systems - Real Time Systems

Mridul Sankar Barik

Jadavpur University

2024

Real-Time Systems

- The operating system, and in particular the scheduler, is the most important component
- Correctness of the system depends not only on the logical result of the computation but also on the time at which the results are produced
- Tasks or processes attempt to control or react to events that take place in the outside world
- These events occur in “real time” and tasks must be able to keep up with them

Soft Real-Time Systems

- In a soft real-time system, it is considered undesirable, but not catastrophic, if deadlines are occasionally missed
- Also known as “best effort” systems
- Most modern operating systems can serve as the base for a soft real time systems
- Examples:
 - Multimedia transmission and reception
 - Networking, Telecom (cellular) networks
 - Web sites and services
 - Computer games

Hard Real-Time Systems

- A hard real-time system has time-critical deadlines that must be met; otherwise a catastrophic system failure can occur
- Requires formal verification/guarantees of being able to always meet its hard deadlines (except for fatal errors)
- Examples:
 - Vehicle Subsystems Control
 - Nuclear Power Plant Control
 - Process Control in Industrial Plants
 - Robotics
 - Air Traffic Control
 - Military Command and Control Systems



Real-Time Tasks

- Hard real-time task
 - One that must meet its deadline
 - Otherwise it will cause unacceptable damage or a fatal error to the system
- Soft real-time task
 - Has an associated deadline that is desirable but not mandatory
 - It still makes sense to schedule and complete the task even if it has passed its deadline


Periodic and Aperiodic Tasks

- Periodic tasks
 - Requirement may be stated as:
 - Once per period T
 - Exactly T units apart
- Aperiodic tasks
 - Has a deadline by which it must finish or start
 - May have a constraint on both start and finish time


Scheduling of Real-Time Process I



Scheduling of Real-Time Process II



(c) Priority-driven preemptive scheduler on preemption points



(d) Immediate preemptive scheduler

Real-Time Scheduling

- Scheduling approaches depend on:
 - Whether a system performs schedulability analysis
 - If it does, whether it is done statically or dynamically
 - Whether the result of the analysis itself produces a scheduler plan according to which tasks are dispatched at run time

Classes of Real-Time Scheduling Algorithms

- Static table-driven approaches
 - Performs a static analysis of feasible schedules of dispatching
 - Result is a schedule that determines, at run time, when a task must begin execution
- Static priority-driven preemptive approaches
 - A static analysis is performed but no schedule is drawn up
 - Analysis is used to assign priorities to tasks so that a traditional priority-driven preemptive scheduler can be used
- Dynamic planning-based approaches
 - Feasibility is determined at run time
 - An arriving task is accepted for execution only if it is feasible to meet its time constraints
- Dynamic best effort approaches
 - No feasibility analysis is performed
 - System tries to meet all deadlines and aborts any started process whose deadline is missed

Deadline Scheduling

- Real-time operating systems are designed with the objective of starting real-time tasks as rapidly as possible and emphasize rapid interrupt handling and task dispatching
- Real-time applications are generally not concerned with sheer speed but rather with completing (or starting) tasks at the most valuable times
- Priorities provide a crude tool and do not capture the requirement of completion (or initiation) at the most valuable time

Information Used for Deadline Scheduling I

- Ready time
 - Time at which task becomes ready for execution
 - For a repetitive or periodic task, this is actually a sequence of times that is known in advance
 - For an aperiodic task, this time may be known in advance, or the operating system may only be aware when the task is actually ready
- Starting deadline
 - Time by which a task must begin
- Completion deadline
 - Time by which task must be completed
 - Typical realtime application will either have starting deadlines or completion deadlines, but not both
- Processing time
 - Time required to execute the task to completion
 - In some cases, this is supplied
 - In others, the OS measures an exponential average

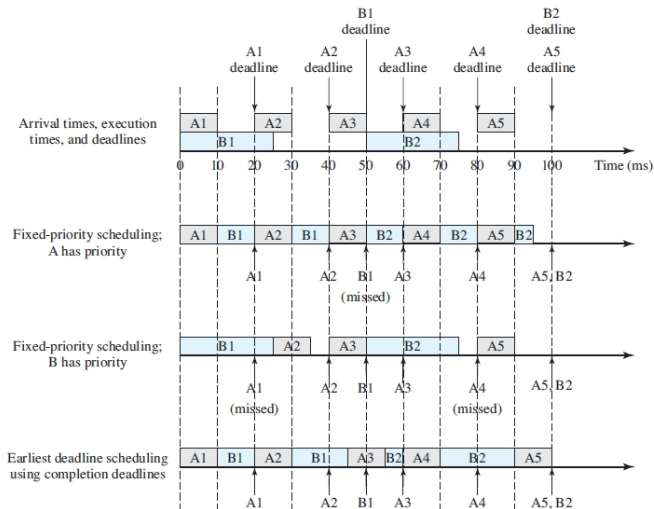
Information Used for Deadline Scheduling II

- Resource requirements
 - Set of resources (other than the processor) required by the task while it is executing
- Priority
 - Measures relative importance of the task
 - Hard real-time tasks may have an “absolute” priority, with the system failing if a deadline is missed
 - If the system is to continue to run no matter what, then both hard and soft realtime tasks may be assigned relative priorities as a guide to the scheduler
- Subtask structure
 - A task may be decomposed into a mandatory subtask and an optional subtask
 - Only the mandatory subtask possesses a hard deadline

Execution Profile of Two Periodic Tasks

Process	Arrival Time	Execution Time	Ending Deadline
A(1)	0	10	20
A(2)	20	10	40
A(3)	40	10	60
A(4)	60	10	80
A(5)	80	10	100
•	•	•	•
•	•	•	•
•	•	•	•
B(1)	0	25	50
B(2)	50	25	100
•	•	•	•
•	•	•	•
•	•	•	•

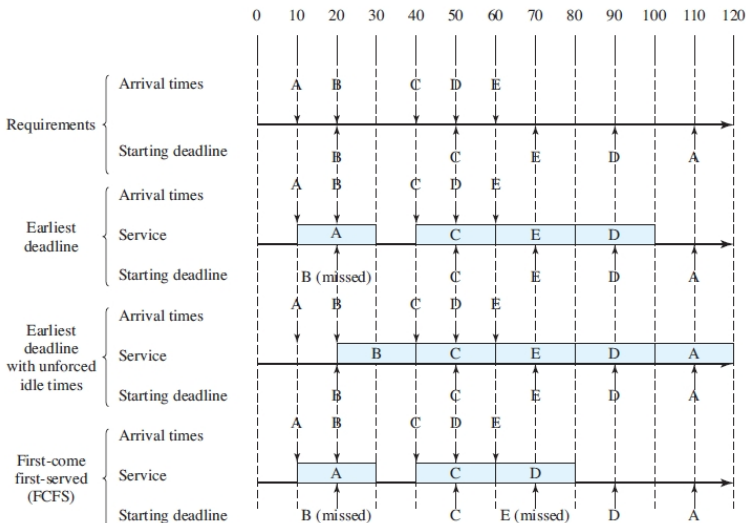
Scheduling of Periodic Real-Time Tasks with Completion Deadlines



Execution Profile of Five Aperiodic Tasks

Process	Arrival Time	Execution Time	Starting Deadline
A	10	20	110
B	20	20	20
C	40	20	50
D	50	20	90
E	60	20	70


Scheduling of Aperiodic Real-Time Tasks with Starting Deadlines




Rate Monotonic Scheduling

- Assigns priorities to tasks on the basis of their periods
 - Highest priority to shortest period

Periodic Task Timing Diagram



Rate Monotonic Scheduling




Periodic Tasks

$T_i = (P_i, C_i)$

P_i = period

C_i = processing time



RMS
Scheduling

Priority Inversion

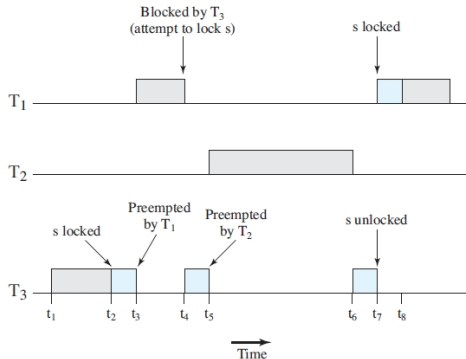
- Can occur in any priority-based preemptive scheduling scheme
- Particularly relevant in the context of real-time scheduling
- Best-known instance involved the Mars Pathfinder mission
- Occurs when circumstances within the system force a higher priority task to wait for a lower priority task
- Unbounded Priority Inversion
 - The duration of a priority inversion depends not only on the time required to handle a shared resource, but also on the unpredictable actions of other unrelated tasks

Mars Pathfinder Mission

- Three tasks with decreasing order of priority
 - T_1 : Periodically checks the health of the spacecraft systems and software
 - T_2 : Processes image data
 - T_3 : Performs an occasional test on equipment status
- T_1 and T_3 share a common data structure protected by a binary semaphore.

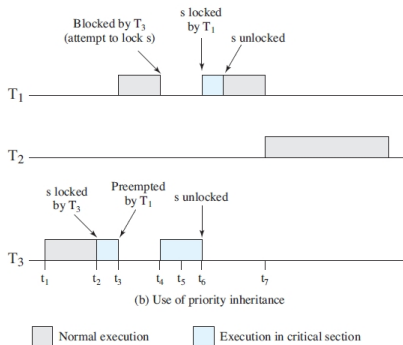
Unbounded Priority Inversion

- t_1 : T_3 begins executing.
- t_2 : T_3 locks semaphore s and enters its critical section.
- t_3 : T_1 , which has a higher priority than T_3 , preempts T_3 and begins executing.
- t_4 : T_1 attempts to enter its critical section but is blocked because the semaphore is locked by T_3 ; T_3 resumes execution in its critical section.
- t_5 : T_2 , which has a higher priority than T_3 , preempts T_3 and begins executing.
- t_6 : T_2 is suspended for some reason unrelated to T_1 and T_3 ; T_3 resumes.
- t_7 : T_3 leaves its critical section and unlocks the semaphore. T_1 preempts T_3 , locks the semaphore and enters its critical section.



Priority Inheritance

- A lower-priority task inherits the priority of any higher-priority task pending on a resource they share
- t_1 : T_3 begins executing.
- t_2 : T_3 locks semaphore s and enters its critical section.
- t_3 : T_1 , which has a higher priority than T_3 , preempts T_3 and begins executing.
- t_4 : T_1 attempts to enter its critical section but is blocked because the semaphore is locked by T_3 . T_3 is immediately and temporarily assigned the same priority as T_1 . T_3 resumes execution in its critical section.
- t_5 : T_2 is ready to execute but, because T_3 now has a higher priority, T_2 is unable to preempt T_3 .
- t_6 : T_3 leaves its critical section and unlocks the semaphore: its priority level is downgraded to its previous default level. T_1 preempts T_3 , locks the semaphore, and enters its critical section.
- t_7 : T_1 is suspended for some reason unrelated to T_2 , and T_2 begins executing.



Priority Ceiling

- A priority is associated with each resource
- The priority assigned to a resource is one level higher than the priority of its highest-priority user
- The scheduler dynamically assigns this priority to any task that accesses the resource
- Once the task finishes with the resource, its priority returns to normal