

Operating Systems - File System Implementation

Mridul Sankar Barik


Jadavpur University

2024

File-System Structure


- File structure
 - Logical storage unit
 - Collection of related information
- File system resides on secondary storage (disks)
 - Provides user interface to storage, mapping logical to physical
 - Provides efficient and convenient access to disk by allowing data to be stored, located retrieved easily
- Disk provides in-place rewrite and random access
 - I/O transfers performed in blocks of sectors (usually 512 bytes)
- File control block – storage structure consisting of information about a file
- Device driver controls the physical device
- File system organized into layers

File System Layers I




- **Device drivers** manage I/O devices at the **I/O control layer**
 - Given commands like “read drive1, cylinder 72, track 2, sector 10, into memory location 1060” outputs low-level hardware specific commands to hardware controller
- **Basic file system:** given command like “retrieve block 123” translates to device driver
 - Also manages memory buffers and caches (allocation, freeing, replacement)
 - Buffers hold data in transit
 - Caches hold frequently used data

File System Layers II



- **File organization module** understands files, logical address, and physical blocks
 - Translates logical block # to physical block #
 - Manages free space, disk allocation
- **Logical file system** manages metadata information
 - Translates file name into file number, file handle, location by maintaining file control blocks (inodes in UNIX)
 - Directory management
 - Protection
- Layering useful for reducing complexity and redundancy, but adds overhead and can decrease performance

File System Layers III




- Many file systems, sometimes many within an operating system
 - Each with its own format
 - CD-ROM is ISO 9660
 - Unix has UFS, FFS
 - Windows has FAT, FAT32, NTFS
 - Linux has more than 40 types, example: ext2 and ext3
 - New ones still arriving – ZFS, GoogleFS, Oracle ASM, FUSE

Allocation Methods

- How to allocate space to files so that disk space is used effectively and files accessed quickly
 - Contiguous Allocation
 - Linked Allocation
 - Indexed Allocation

Contiguous Allocation




directory

file	start	length
count	0	2
tr	14	3
mail	19	6
list	28	4
f	6	2


- Each file occupies a set of contiguous blocks on the disk
- Number of disk seeks required for accessing contiguously allocated files is minimal
- Directory entry contains starting location (block #) and length (number of blocks)
- Both sequential and direct access is supported
- Application of dynamic storage-allocation (how to satisfy a request of size n from a list of free holes)
- External fragmentation may occur
 - Solution: Defragmentation
- Files cannot grow
 - Solution: Allocate in contiguous chunks
 - Location of file's blocks are recorded as Location, Block count, Link to the first block of the next chunk

Linked Allocation




- Each file is a linked list of disk blocks, blocks may be scattered anywhere on the disk
- Directory entry contains a pointer to the first and last blocks of the file
- No external fragmentation
- File size need not be declared during creation
- Can be use only for sequential access file
- Space required for pointers are wasted
- Solutions
 - Allocate space in clusters (ex - 4 blocks)
 - Increases internal fragmentation
- Reliability
 - A bug in OS software or disk hardware failure may cause pointer to be lost or damaged
 - Use doubly linked list
 - Store name of file, relative block # in each block

File Allocation Table



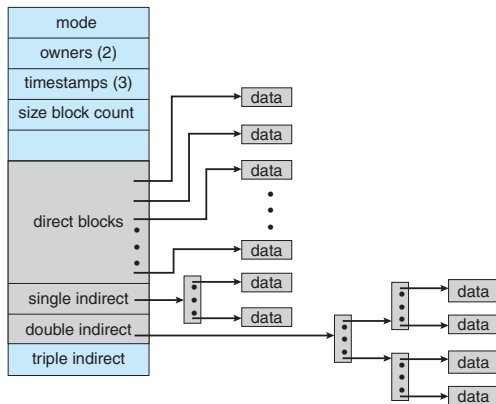
- A section of disk at the beginning of each partition is set aside to contain FAT
- FAT contain
 - One entry for each block; indexed by block number
 - Directory entry contain block number of the first block of the file
 - Last block of a file contain special end of file marker
 - Unused blocks in FAT contain 0 value

Indexed Allocation I



- Brings all pointers together into the index block
- Directory entry contains the address of the index block
- All pointers in the index block are NULL initially
- i^{th} entry in the index block points to the i^{th} block within the file
- Supports direct access
- Pointer overhead is greater than linked allocation

Indexed Allocation II



- How large the index block should be?
- Linked scheme:
 - An index block is normally one disk block
 - For large files, link together several index blocks
- Multilevel index:
 - First level index block points to a set of second level index blocks which contain pointers to the file blocks
- Combined Scheme:
 - Ex: Unix File System