

Operating Systems - Main Memory

Mridul Sankar Barik

Jadavpur University


2024

Background

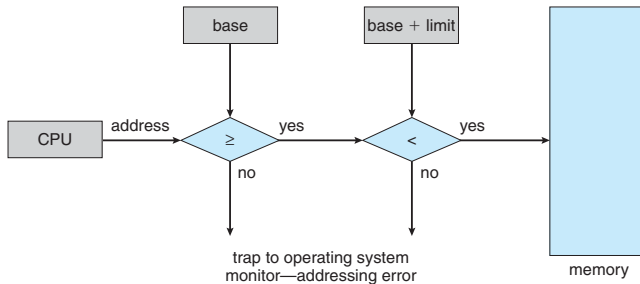
- Memory is large array of words or bytes each with its own address
- Program must be brought (from disk) into memory and placed within a process for it to be run
- Main memory and registers are only storage CPU can access directly
- CPU fetches instructions from memory according to the value of register Program Counter (PC)
- These instructions may cause additional loading from or storing to specific memory addresses
- Memory unit sees only a stream of memory addresses; it does not know
 - How they are generated
 - What they are for (instruction or data)
- Cache sits between main memory and CPU registers
- Protection of memory required to ensure correct operation

Base and Limit Registers

- A pair of base and limit registers define the logical address space
- CPU must check every memory access generated in user mode to be sure it is between base and limit for that user



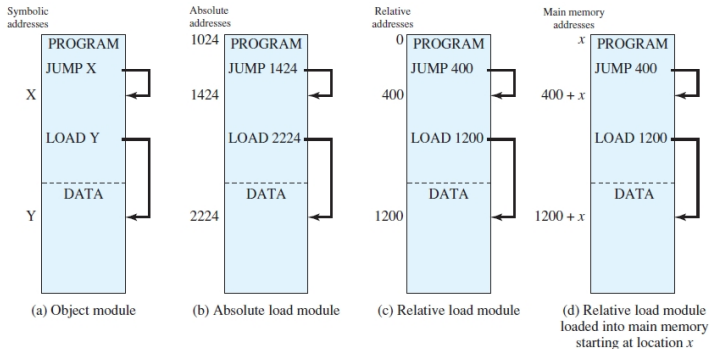
Hardware Address Protection



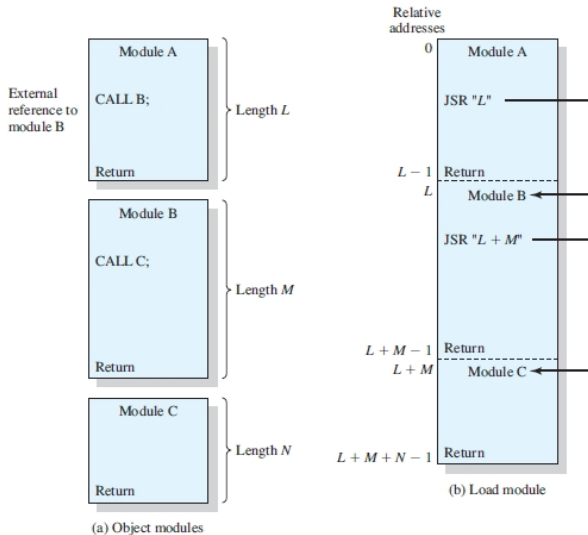
Address Binding

- Addresses are represented in different ways at different stages of a program's life
 - Source code addresses usually symbolic
 - Compiled code addresses bind to relocatable addresses, i.e. "14 bytes from beginning of this module"
 - Linker or loader will bind relocatable addresses to absolute addresses, i.e. 74014
 - Each binding maps one address space to another
- Address binding of instructions and data to memory addresses can happen at three different stages
 - **Compile time:** If memory location known a priori, absolute code can be generated;
 - Must recompile code if starting location changes
 - Example – MSDOS.COM programs
 - **Load time:** Compiler must generate relocatable code if memory location is not known at compile time
 - **Execution time:** Binding delayed until run time if the process can be moved during its execution from one memory segment to another
 - Need hardware support for address maps (e.g., base and limit registers)

Absolute and Relocatable Load Module



The Linking Function



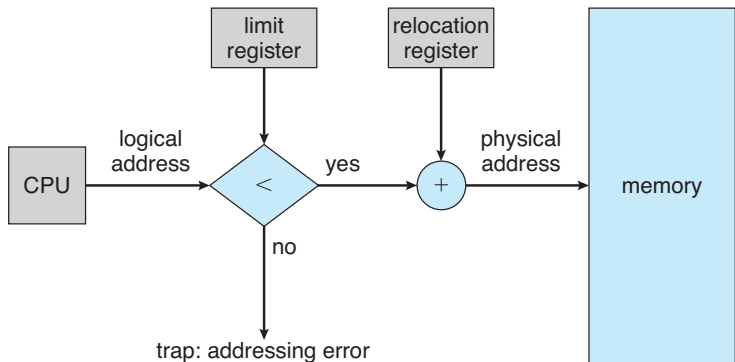
Logical vs. Physical Address Space

- The concept of a logical address space that is bound to a separate physical address space is central to proper memory management
 - **Logical address** – generated by the CPU; also referred to as virtual address.
 - **Physical address** – address seen by the memory unit.
- Logical and physical addresses are the same in compile-time and load-time address-binding schemes
- Logical (virtual) and physical addresses differ in execution-time address-binding scheme
- Memory Management Unit
 - Hardware device that maps virtual to physical address
 - The user program deals with logical addresses; it never sees the real physical addresses
- **Logical address space** is the set of all logical addresses generated by a program
- **Physical address space** is the set of all physical addresses generated by a program

Contiguous Memory Allocation

- Main memory is usually divided into two partitions
 - Resident operating system, usually held in low memory with interrupt vector
 - User processes held in high memory
- Each process is contained in a single contiguous section of memory
 - Relocation-register scheme used to protect user processes from each other, and from changing operating-system code and data
 - Relocation register contains value of smallest physical address; limit register contains range of logical addresses – each logical address must be less than the limit register

Hardware Support for Relocation and Limit Registers



Multiple Partition Allocation

- Multiple-partition allocation
 - Degree of multiprogramming is limited by number of partitions
 - Variable-partition sizes used for efficiency
 - Hole – block of available memory; holes of various size are scattered throughout memory
 - When a process arrives, it is allocated memory from a hole large enough to accommodate it
 - Process exiting frees its partition, adjacent free partitions are combined
 - Operating system maintains information about: (a) allocated partitions, (b) free partitions (hole)

Dynamic Storage-Allocation Problem

- How to satisfy a request of size n from a list of free holes?
 - First-fit: Allocate the first hole that is big enough
 - Best-fit: Allocate the smallest hole that is big enough; must search entire list, unless ordered by size
 - Produces the smallest leftover hole
 - Worst-fit: Allocate the largest hole; must also search entire list
 - Produces the largest leftover hole

Fragmentation

- External fragmentation – total memory space exists to satisfy a request, but it is not contiguous
- Internal fragmentation – allocated memory may be slightly larger than requested memory; this size difference is memory internal to a partition, but not being used
- Reduce external fragmentation by
 - 1 Compaction
 - Shuffle memory contents to place all free memory together in one large block
 - Compaction is possible only if relocation is dynamic, and is done at execution time
 - 2 Allow logical address space of a process to be non contiguous
 - Paging
 - Segmentation

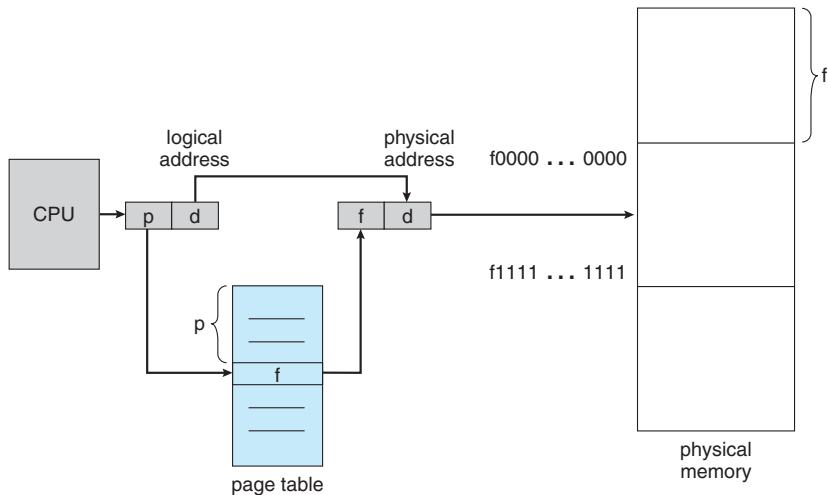
Paging

- Physical address space of a process can be noncontiguous; process is allocated physical memory whenever the latter is available
 - Avoids external fragmentation
 - Avoids problem of varying sized memory chunks
- Divide physical memory into fixed-sized blocks called **frames**
 - Size is power of 2, between 512 bytes and 16 Mbytes
- Divide logical memory into blocks of same size called **pages**
- Keep track of all free frames
- To run a program of size N pages, OS needs to find N free frames and load program
- Set up a page table to translate logical to physical addresses

Address Translation Scheme

- Address generated by CPU is divided into:
 - Page number (p) – used as an index into a page table which contains base address of each page in physical memory
 - Page offset (d) – combined with base address to define the physical memory address that is sent to the memory unit

Paging Hardware



Paging Example

0	a
1	b
2	c
3	d
4	e
5	f
6	g
7	h
8	i
9	j
10	k
11	l
12	m
13	n
14	o
15	p

logical memory

0	5
1	6
2	1
3	2

page table

0	
4	i j k l
8	m n o p
12	
16	
20	a b c d
24	e f g h
28	

physical memory

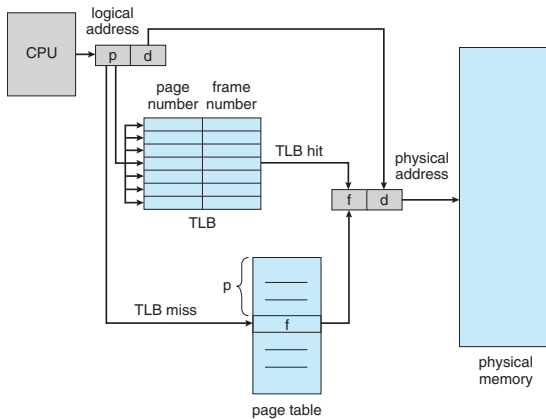
Paging

- Calculating internal fragmentation
 - Page size = 2,048 bytes
 - Process size = 72,766 bytes
 - 35 pages + 1,086 bytes
 - Internal fragmentation of $2,048 - 1,086 = 962$ bytes
 - Worst case fragmentation = 1 frame – 1 byte
 - On average fragmentation = $1 / 2$ frame size
 - So small frame sizes desirable?
 - But each page table entry takes memory to track
- Process view and physical memory now very different
- By implementation process can only access its own memory

Implementation of Page Table

- Page table is kept in main memory
- Page-table base register (PTBR) points to the page table
- Page-table length register (PTLR) indicates size of the page table
- In this scheme every data/instruction access requires two memory accesses
 - One for the page table and one for the data / instruction
- The two memory access problem can be solved by the use of a special fast-lookup hardware cache called associative memory or translation look-aside buffers (TLBs)
- TLBs typically small (64 to 1,024 entries)
- On a TLB miss, entry is loaded into the TLB for faster access next time
 - Replacement policies must be considered

Paging Hardware With TLB



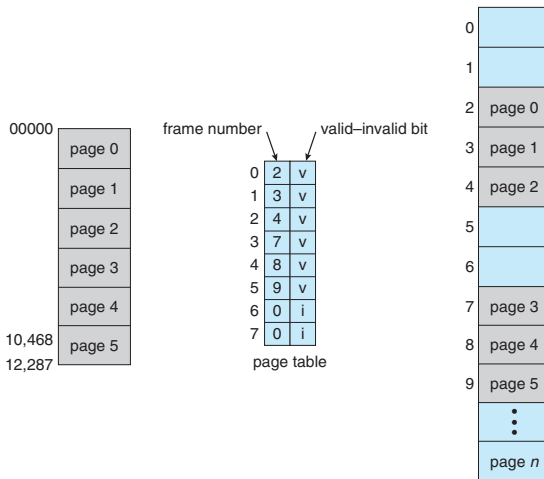
Effective Access Time

- Associative Lookup = ϵ time unit
 - Can be $< 10\%$ of memory access time
- Hit ratio = α
 - Hit ratio – percentage of times that a page number is found in the associative registers; ratio related to number of associative registers
- Consider $\alpha = 80\%$, $\epsilon = 20$ ns for TLB search, $m_a = 100$ ns for memory access
- Effective Access Time (EAT)
 - $EAT = (m_a + \epsilon)\alpha + (2m_a + \epsilon)(1-\alpha)$
- $EAT = 0.80 \times (100 + 20) + 0.20 \times (200 + 20) = 96 + 44 = 140$ ns
- Consider more realistic hit ratio $\Rightarrow \alpha = 99\%$, $\epsilon = 20$ ns for TLB search, $m_a = 100$ ns for memory access
 - $EAT = 0.99 \times (100 + 20) + 0.01 \times (200 + 20) = 118.8 + 2.2 = 121$ ns

Memory Protection

- Memory protection implemented by associating protection bit with each frame to indicate if read-only or read-write access is allowed
 - Can also add more bits to indicate page execute-only, and so on
- Valid-invalid bit attached to each entry in the page table:
 - “valid” indicates that the associated page is in the process’ logical address space, and is thus a legal page
 - “invalid” indicates that the page is not in the process’ logical address space
- Any violations result in a trap to the kernel

Valid (v) or Invalid (i) Bit In A Page Table




Structure of the Page Table

- Memory structures for paging can get huge using straight-forward methods
 - Consider a 32-bit logical address space as on modern computers
 - Page size of 4 KB (2^{12})
 - Page table would have ($2^{32}/2^{12} = 2^{20}$) entries
 - If each entry is 4 bytes \rightarrow 4 MB of physical address space / memory for page table alone
 - That amount of memory may not be available in contiguous areas in main memory
- Hierarchical Paging
- Hashed Page Tables
- Inverted Page Tables

Hierarchical Page Tables

- Break up the logical address space into multiple page tables
- A simple technique is a two-level page table
- We then page the page table


Two-Level Page-Table Scheme



Two-Level Paging Example

- A logical address (on 32-bit machine with 2K page size) is divided into:
 - a page number consisting of 20 bits
 - a page offset consisting of 12 bits
- Since the page table is paged, the page number is further divided into:
 - a 10-bit page number
 - a 10-bit page offset


Address-Translation Scheme



Hashed Page Tables I

- Common in address spaces > 32 bits
- The virtual page number is hashed into a page table
 - This page table contains a chain of elements hashing to the same location
- Each element contains (1) the virtual page number (2) the value of the mapped page frame (3) a pointer to the next element
- Virtual page numbers are compared in this chain searching for a match
 - If a match is found, the corresponding physical frame is extracted

Hashed Page Tables II



Inverted Page Table I

- Rather than each process having a page table and keeping track of all possible logical pages, track all physical frames
- One entry for each real frame of memory
- Entry consists of the virtual address of the page stored in that real memory location, with information about the process that owns that page
- Decreases memory needed to store each page table, but increases time needed to search the table when a page reference occurs

Inverted Page Table II

